

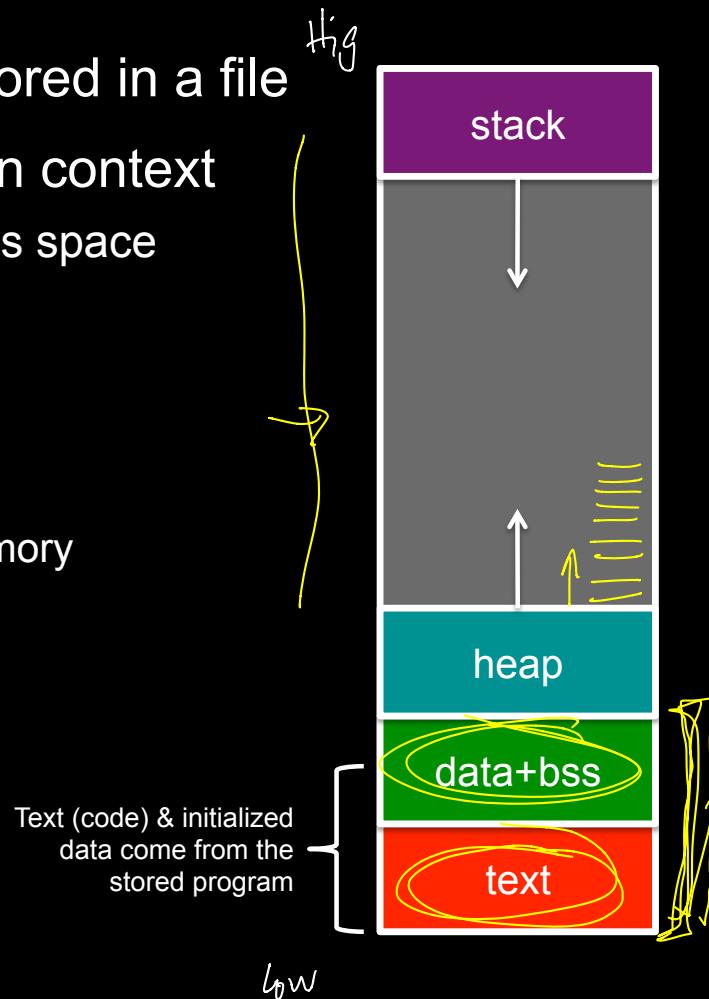
# Operating Systems Design

## 4. Processes

Paul Krzyzanowski  
[pxk@cs.rutgers.edu](mailto:pxk@cs.rutgers.edu)

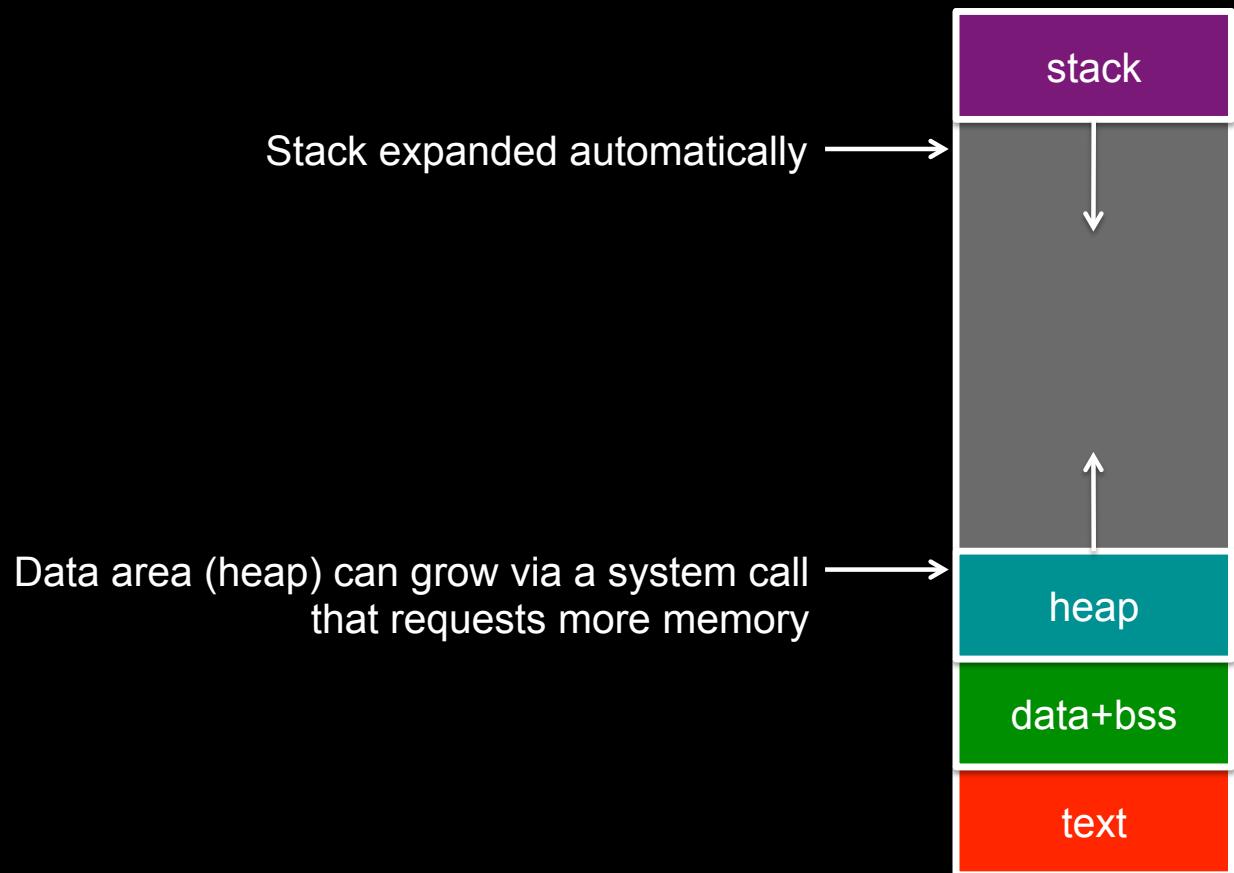
# Process

- Program: code & static data stored in a file
- Process: a program's execution context
  - Each process has its own address space
  - Memory map *Mapping*
    - Text: compiled program
    - Data: initialized static data
    - BSS: uninitialized static data
    - Heap: dynamically allocated memory
    - Stack: call stack
  - Process context:
    - Program counter
    - CPU registers



# Growing memory

---



# Contexts

---

- Entering the kernel
  - Hardware interrupts
    - Asynchronous events (I/O, clock, etc.)
    - Do not relate to the context of the current process
      - Because they are asynchronous, *any* process might be running when they occur
  - Software traps
    - Are related to the context of the current process [process context]
    - Examples: illegal memory access, divide by zero, illegal instruction
  - Software initiated traps
    - System call from the current process [process context]
    - *The current executing process' address space is active on a trap*
- Saving state
  - Kernel stack switched in upon entering kernel mode
  - Kernel must save machine state before servicing event
    - Registers, flags (program status word), program counter, ...

*read*

# System calls

---

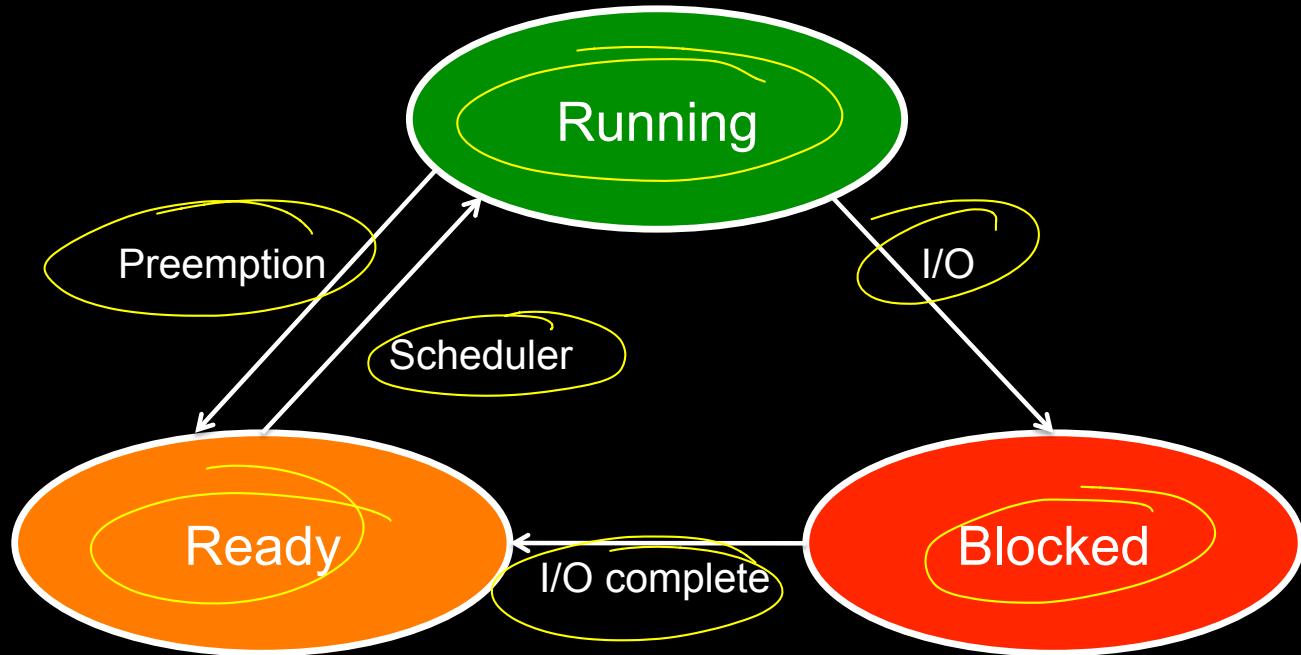
- Entry: Trap to system call handler
  - Save state
  - Verify parameters are in a valid address
  - Copy them to kernel address space
  - Call the function that implements the system call
    - If the function has to (cannot be satisfied immediately) then
      - Context switch to let another ready process run
      - Put our process on a blocked list
- Return from system call or interrupt
  - Check for signals to the process
    - Call the appropriate handler if signal is not ignored
  - Check if another process should run
    - Context switch to let the other process run
    - Put our process on a *ready* list
  - Calculate time spent in the call for profiling/accounting
  - Restore user process state
  - Return from interrupt

# Processes in a Multitasking Environment

---

- Multiple concurrent processes
  - Each has a unique identifier: **Process ID (PID)**
- Asynchronous events (interrupts) may occur
- Processes may request operations that take a long time
- Goal: **have some process running at all times**
- Context saving/switching
  - Processes may be suspended and resumed
  - Need to save all state about a process so we can restore it

# Process states



CPU-bound process → (Almost) no I/O

# Keeping track of processes

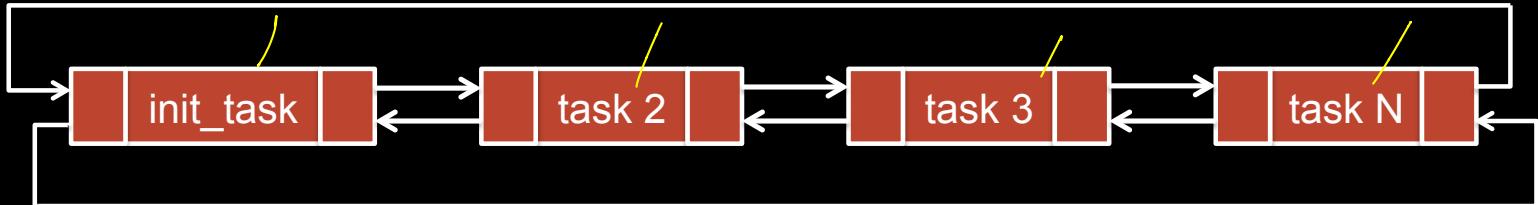
---

- Process list stores a **Process Control Block (PCB)** per process
- A Process Control Block contains:
  - Process ID
  - Machine state (registers, program counter, stack pointer)
  - Parent & list of children
  - Process state (ready, running, blocked)
  - Memory map
  - Open file descriptors
  - Owner (user ID) – determine access & signaling privileges
  - Event descriptor if the process is blocked
  - Signals that have not yet been handled
  - Policy items: Scheduling parameters, memory limits
  - Timers for accounting (time & resource utilization)
  - (Process group)

# Processes in Linux

- The OS creates one task on startup:  
*init*: the parent of all tasks  
*launchd*: replacement for *init* on Mac OS X and FreeBSD
- Process state stored in struct `task_struct`
  - Defined in `linux/sched.h`
- Stored as a circular, doubly linked list
  - `struct list_head` in `linux/list.h` 

```
struct task_struct init_task; /* static definition */
```



Aside:

## The Linux kernel

- ① Download from: `kernel.org` (latest stable)
- ② - Unzip file: `unxz filename.tar.xz`
- ③ - Extract files: `tar xf filename.tar`
- ④ - Open <sup>any</sup> file:  
e.g. `vi include/linux/sched.h`

(We will build the kernel soon, insha'Allah)

grep → " " .  
whole tree      any string      dot (starting location)

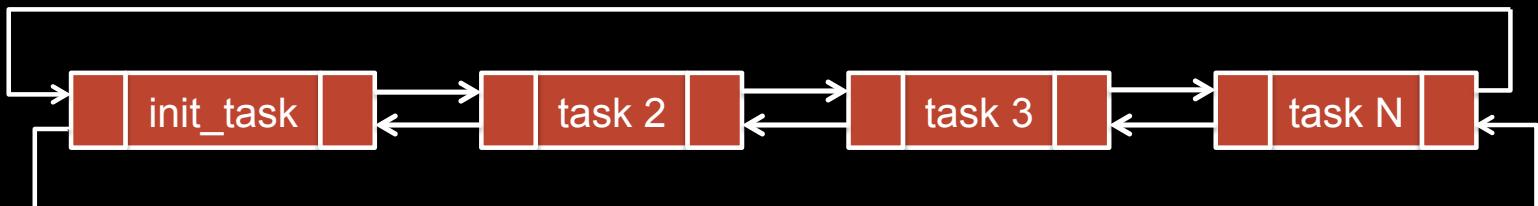
# Processes in Linux

- Iterating through processes

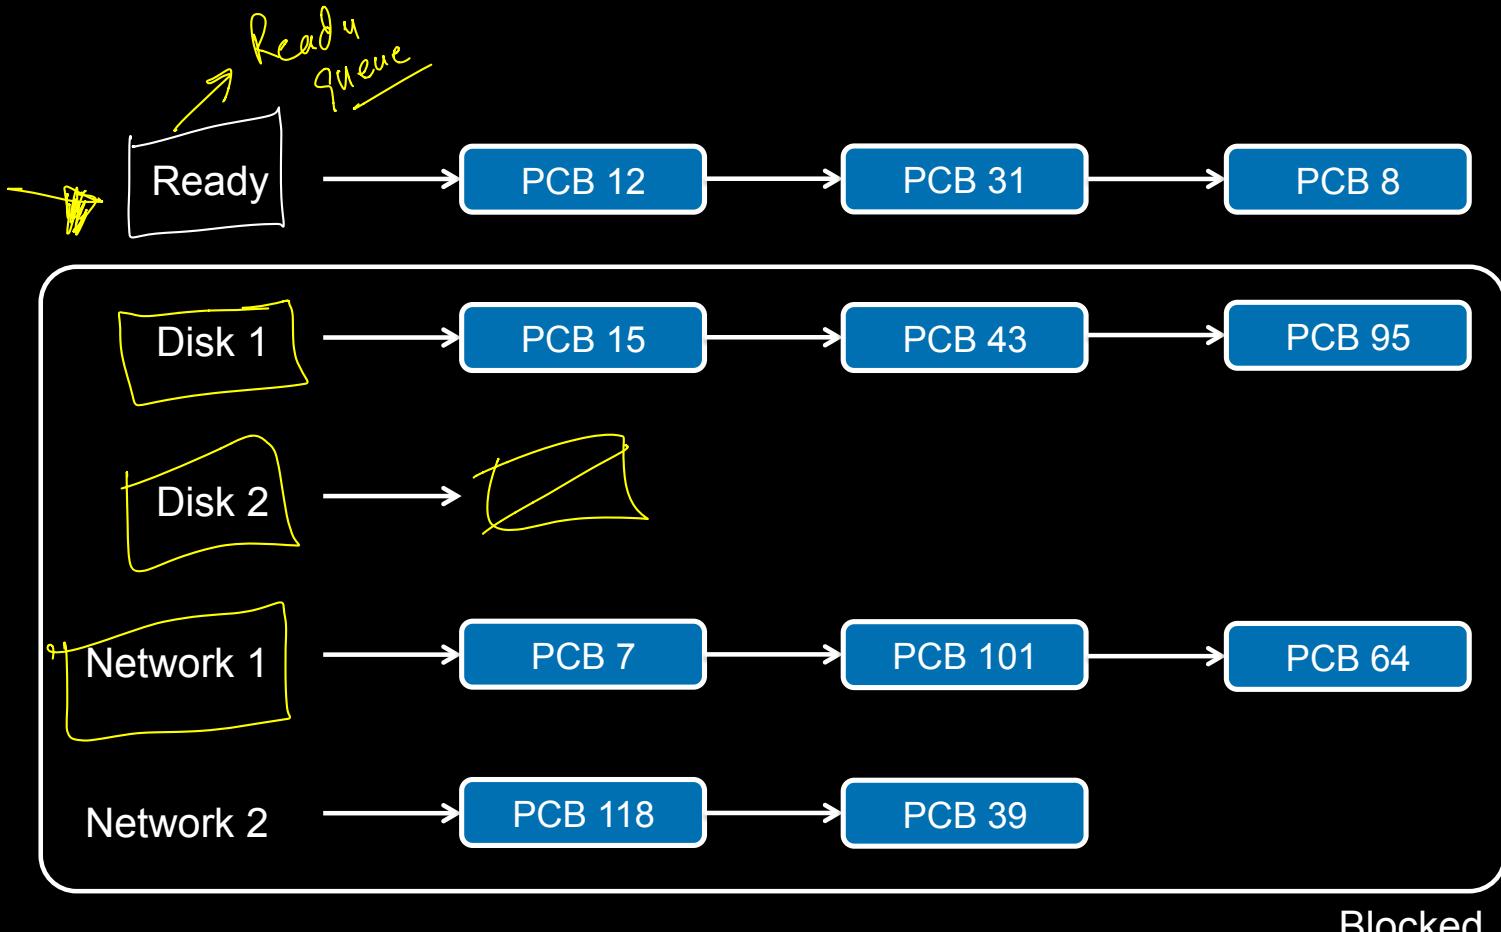
```
for (p = &init_task; ((p = next_task(p)) != &init_task; ) {  
    /* whatever */  
}
```

- The current process on the current CPU is obtained from the macro **current**

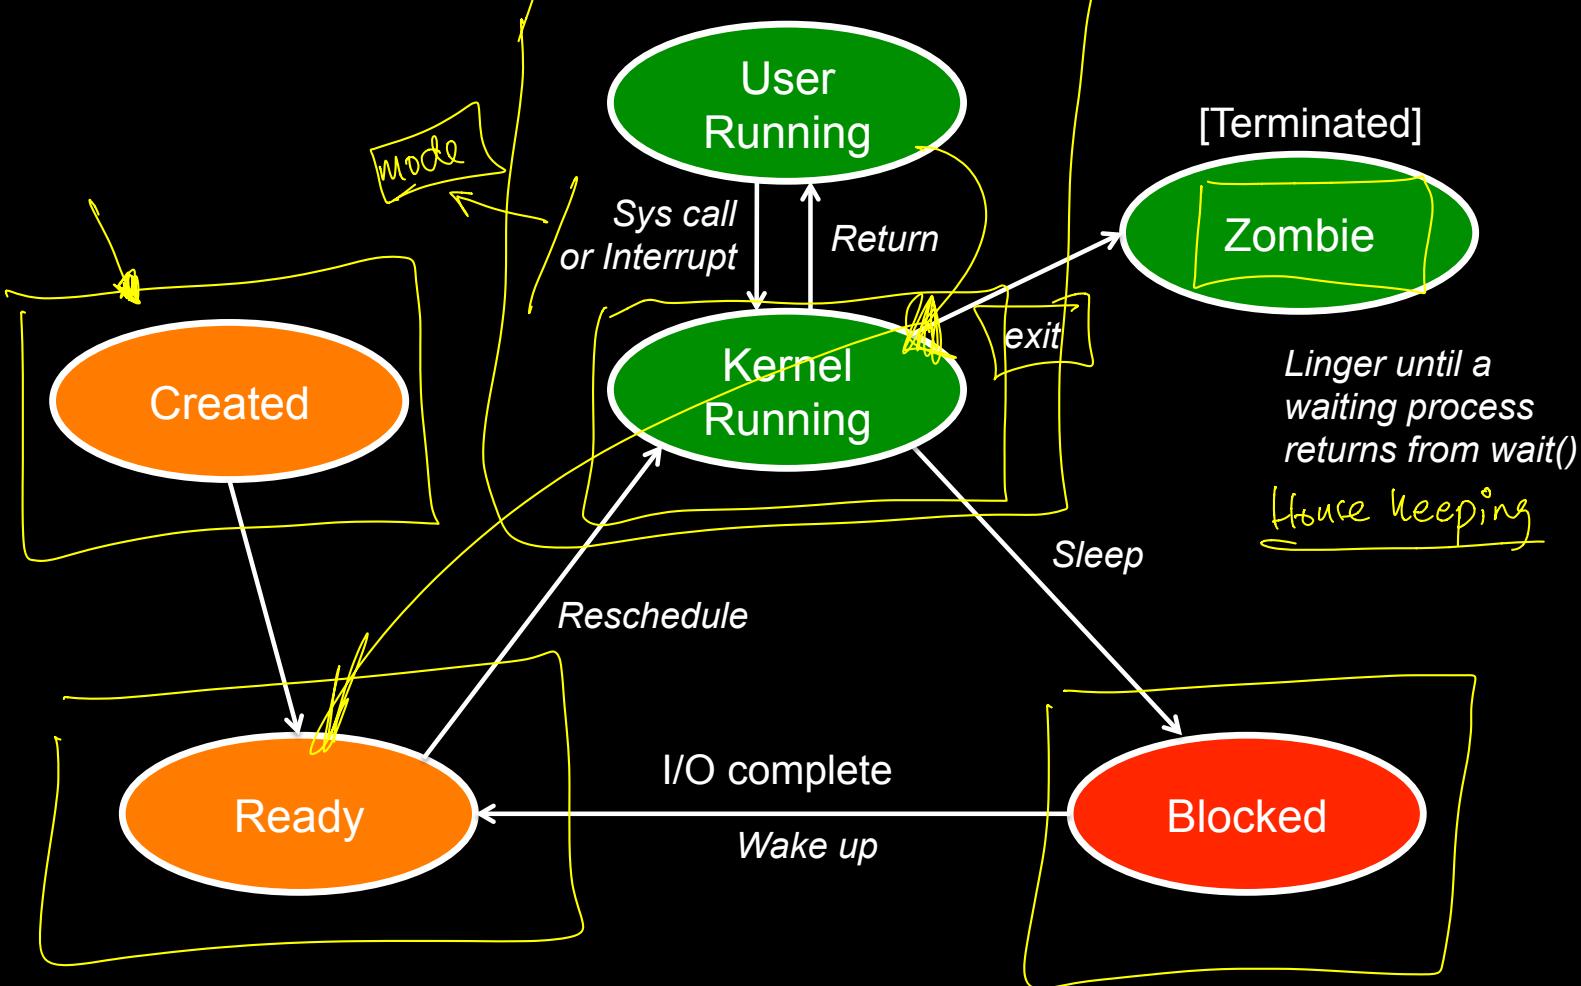
```
current->state = TASK_STOPPED;
```



# Processes on Ready & Blocked Queues



# Process States: a bit more detail



# Creating a process under POSIX

## **fork** system call

- Clones a process into two processes
  - New context is created: duplicate of parent process
- *fork* returns 0 to the child and the process ID to the parent

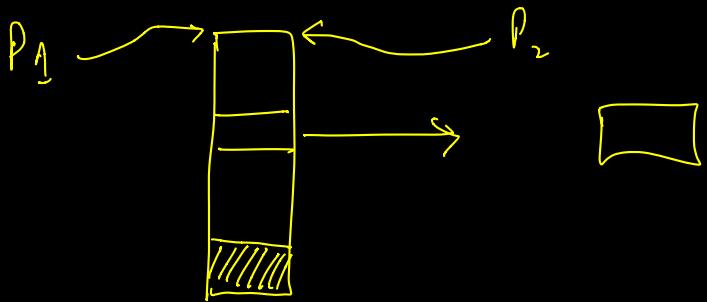
*"fork in  
the road"*



# What happens?

fork()

- Check for available resources
- Allocate a new PCB
- Assign a unique PID
- Check process limits for user
- Set child state to “created”
- Copy data from parent PCB slot to child
- Increment counts on current directory & open files
- Copy parent context in memory (or set *copy on write*)
- Set child state to “ready to run”
- Wait for the scheduler to run the process



→ Read → no copy  
 Write → "Buri baat"

(Create a) "Copy on write"  
CoW

# Fork Example

---

```
#include <stdio.h>

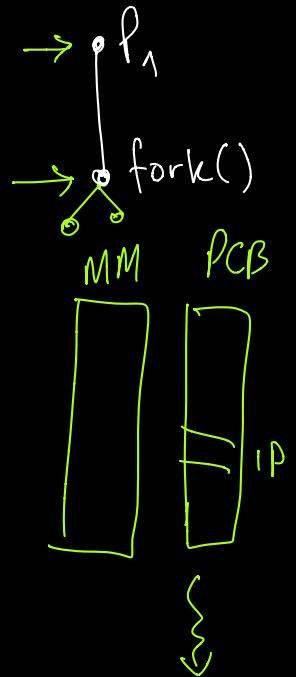
main(int argc, char **argv) {
    int pid;

    switch (pid=fork()) {
        case 0:   printf("I'm the child\n");
                   break;
        default:
            printf("I'm the parent of %d\n", pid);
            break;
        case -1:
            perror("fork");
    }
}
```

```

int main( ... ) {
    int pid;
    pid = fork();
    if (pid == 0)
        printf("child");
}

```



```

else if (pid == -1)
    printf("Error");
}

```

```

else
    printf("Parent of %d", pid);
}

```

$P_i \rightarrow \text{Child}$

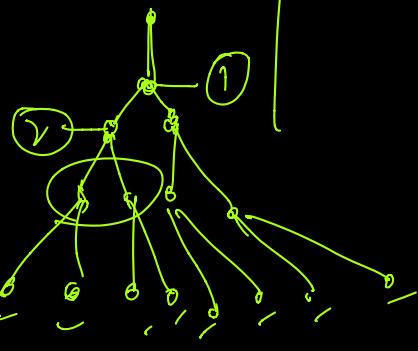
$P_i \rightarrow \text{Parent of } \_\_$

Parent of  $\_\_$   
Child.

```

int main() {
    fork();
    fork();
    fork();
}

```



main()

int pid = fork();

pid = fork();

if (pid == 0)

fork() ←

fork() ←

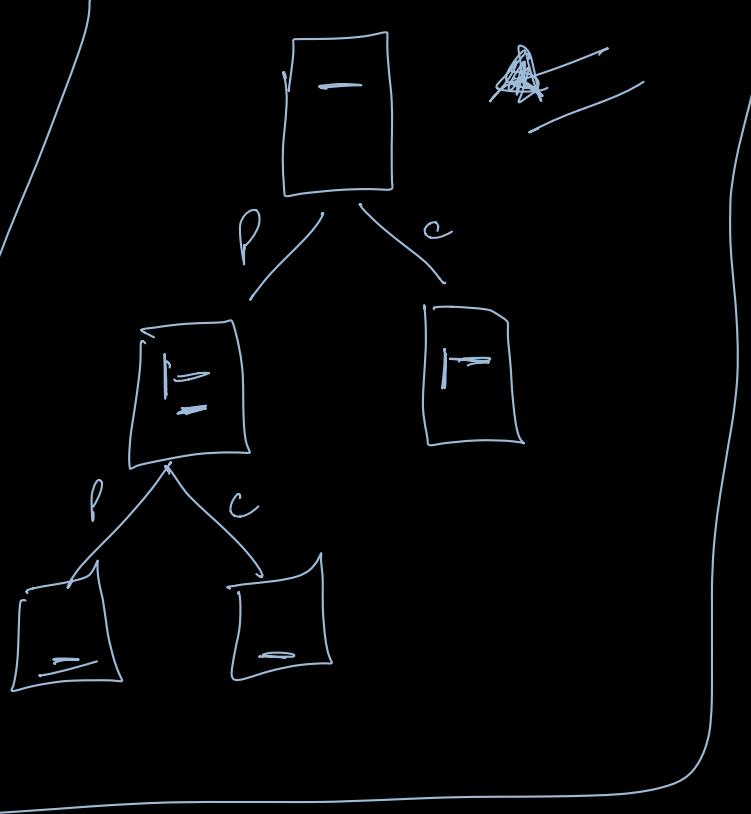
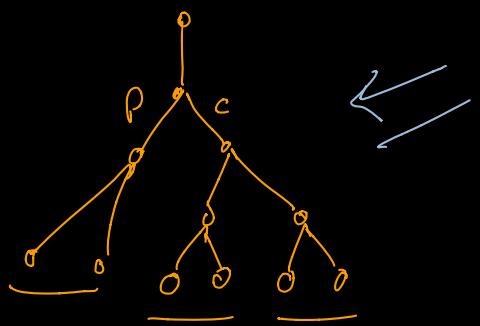
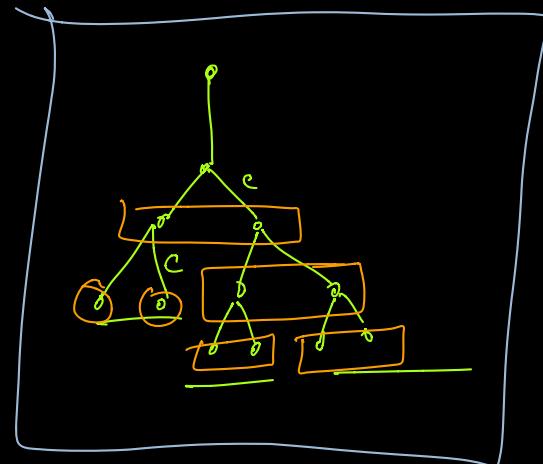
else

fork() ↑

i = 0

while (i < 4)

fork() →  
i++



fork  
for  
for

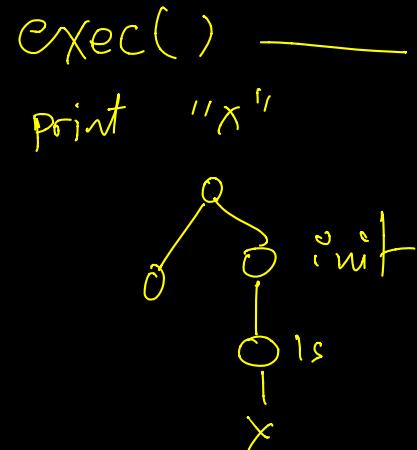
# Running other programs

**execve**: replace the current process image with a new one

- See also exec, execle, execlp, execvp, execvP

- New program inherits:

- Processes group ID ✓
- Open files ✓
- Access groups
- Working directory ✓
- Root directory ✓
- Resource usages & limits ✓
- Timers ✓
- File mode mask
- Signal mask



# Exec Example

```
#include <unistd.h>
```

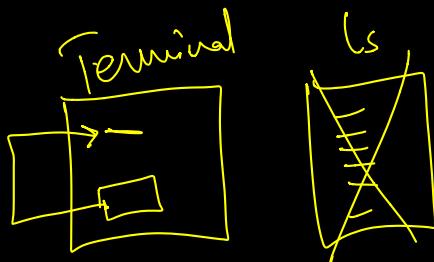
```
main(int argc, char **argv) {
    char *av[] = { "ls", "-al", "/", 0 };
    execvp("ls", av);
}
```

printf "x"



Parameter.

Char \*x → string



Read

Eval

Print

loop

fork  
+  
exec

x

# Fork & exec combined

---

- UNIX creates processes via *fork* followed by *exec*
- Windows approach
  - *CreateProcess* system call to create a new child process
  - Specify the executable file and parameters
  - Identify startup properties (windows size, input/output handles)
  - Specify directory, environment, and whether open files are inherited

# Exiting a process

---

***exit*** system call

```
#include <stdlib.h>

main(int argc, char **argv) {
    exit(0);
}
```



# exit: what happens?

- Ignore all signals
- If the process is associated with a controlling terminal
  - Send a hang-up signal to all members of the process group
  - reset process group for all members to 0
- close all open files
- release current directory
- release current changed root, if any
- free memory associated with the process
- write an accounting record (if accounting)
- make the process state zombie
- assign the parent process ID of any children to be 1 (init)
- send a “death of child” signal to parent process (SIGCHLD)
- context switch (we have to!)

# Wait for a child process to die

---

## *wait* system call

- Suspend execution until a child process exits
- *wait* returns the exit status of that child.

```
int pid, my_pid, status;

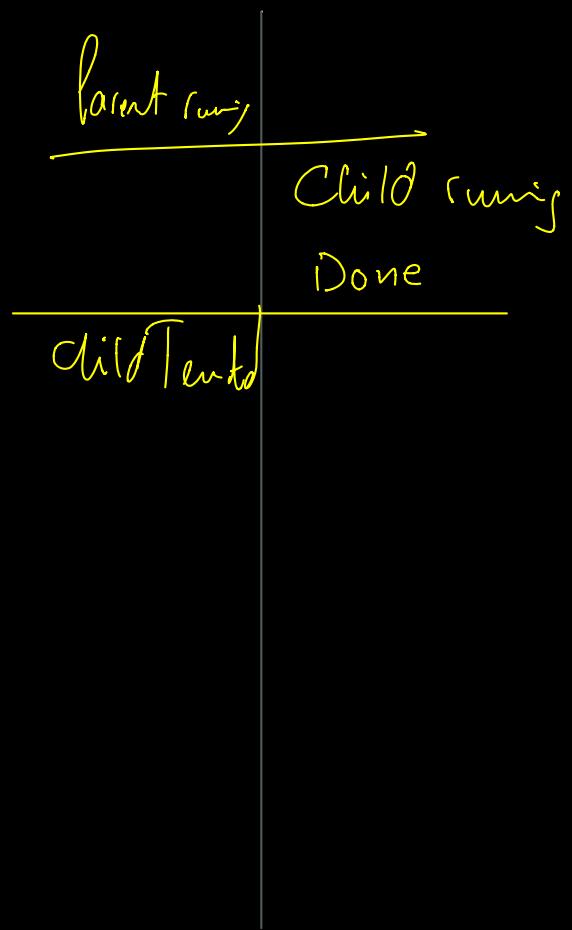
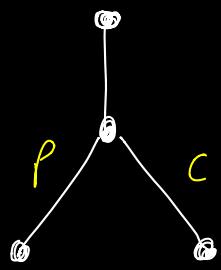
switch (my_pid=fork()) {
case 0:          /* do child stuff */ break;
case -1:         /* do error stuff */ break;

default:         /* wait for child to exit */
    while (pid=wait(&status))
        if (pid==my_pid)
            printf("got exit of %d\n", WEXITSTATUS(status));
        break;
}
```

```

int main () {
    int pid;
    pid = fork();
    if (pid == 0) { // Child running
        "Child running" || ls
    } else { // Parent running
        "Parent running"
        wait (NULL);
        "Child terminated"
    }
}

```



Inputs

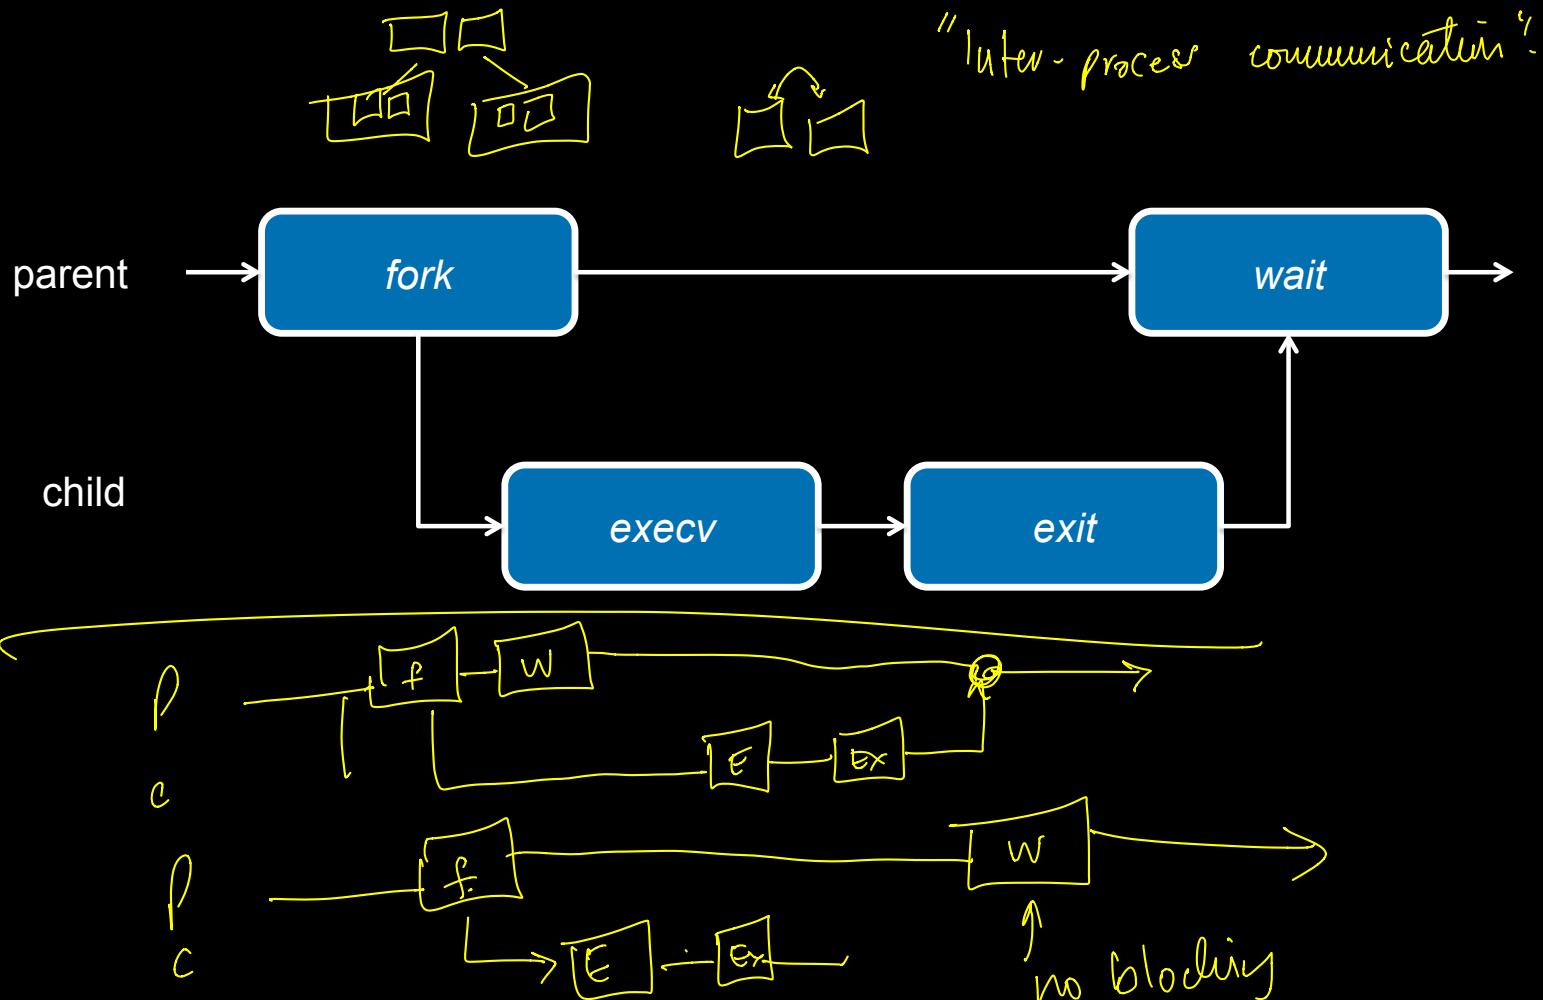
```

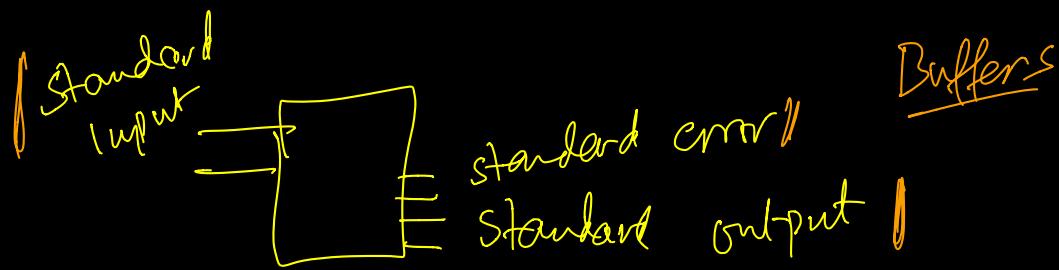
fork
parent: — wait
child → exec
↓
exit
parent: —

```

IN  
↳ GUI login  
↳ Desktop  
↳ T

# Parent & child processes

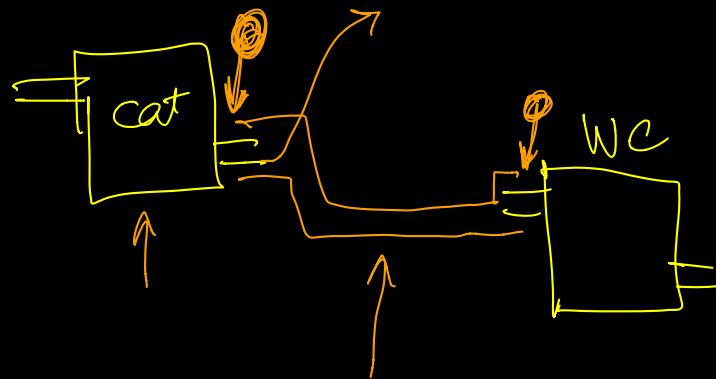




Buffers

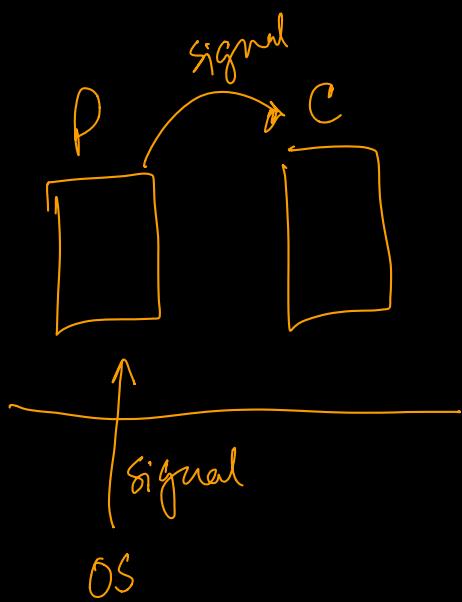
standard err //

standard output |



Inter-process Communication

(IPC)



cascading



# Signals

---

- Inform processes of asynchronous events
  - Processes may specify signal handlers
- Processes can poke each other (if they are owned by the same user)

- Sending a signal:
  - `kill (int pid, int signal_number)`
- Detecting a signal:
  - `signal (signal_number, function)`

*Kill → send a signal*

*signal → receive a signal*

*Kill [pid]*

<signal.h>

```
int main() {  
    signal(SIGINT, handler);
```

while (1);

}

"event driven  
programming"

```
"register"  
void handler (int sig) {  
  
    "signal", sig;
```

}

# The End

$P_1 \rightarrow$  fork

$P_1$

$P_1 \rightarrow$

$P_2 \rightarrow$

(000)000

sum



$P_1 \rightarrow P_2$

MM  $\rightarrow$  not duplicated -(COW)

|| context switch  $\rightarrow$  One core  $\times$   
# duplicates  $>$  # cores

"Light-weight processes"